

Evidence Before Improvement: A White-Box Contract Audit for Autonomous MLIP Research

Anonymous Authors¹

Abstract

Autonomous research agents can optimize a reported metric while silently losing the evidence needed to justify the result. We preregister a deterministic white-box contract-coverage audit of an artifact-bound research trace for machine-learning interatomic potentials (MLIPs). Twelve hand-authored fault operators were each applied to four trace instantiations, yielding 48 single-fault cases and four clean controls. The semantic grader detected 48/48 declared corruptions with no clean rejection; a mutable manifest-only integrity ablation detected 4/48. A separate, previously verified two-iteration MACE replay on an NVIDIA L4 illustrates the policy consequence: one candidate crossed the frozen metric-improvement threshold but remained rejected after its exact-equality reproducibility gate failed. These observations support an infrastructure claim about coverage of the tested contract and fault catalog, not general grader robustness, improved MACE performance, or scientific validity.

1. Introduction

Language-model agents can plan, execute, and write machine-learning experiments (Huang et al., 2024; Yamada et al., 2025), while recent benchmarks evaluate whether agents can reproduce research artifacts (Starace et al., 2025). Success metrics alone, however, do not establish that an experimental conclusion is reproducible: reproducible code and reproducible findings are distinct goals (Bouthillier et al., 2019). This distinction is acute for MLIPs, where model identity, energy and force units, data partitions, hardware, and validation protocols all affect interpretation (Battaglia et al., 2022; Maxson et al., 2024).

We study a narrower question: *does semantic re-derivation*

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.

detect repository-defined trace faults more completely than producer-declared hashes? Drawing on fault injection as a controlled dependability-analysis technique (Hsueh et al., 1997), our contribution is a preregistered white-box contract audit over 12 fault families and four seeds, paired with a real-MACE negative-result case study. The paper intentionally stops at an infrastructure claim ceiling.

2. Method and Agent Workflow

Evidence contract. The controller emits sealed objective, proposal, legality, execution, evaluation, decision, lesson, lineage, and manifest artifacts. The evaluator has a distinct code identity, and the decision is recomputed from frozen acceptance constraints. Reflection and reviewer text may influence a later proposal but cannot support a numerical claim. This follows the broader principle of falsification-oriented automated analysis (Huang et al., 2025), but our grader is deterministic and does not use an LLM.

3. Experimental Protocol

Preregistered comparison. For each seed from 20260712–20260715, we generated one clean two-iteration synthetic trace and then made twelve separate derived copies. Each copy received exactly one frozen fault: missing Tea Time, missing lesson, evaluator identity collision, immutable-input change, forbidden label payload, unauthorized remote attestation, decision–metric inconsistency, internal-seal tamper, registered-artifact corruption, inflated scientific status, forged rejection without rollback, or broken lineage parent. The recorded execution evaluated every seed–fault pair once and reports `retry_count=0`. The integrity ablation verifies every path, size, and SHA-256 declared in a mutable producer manifest. For semantic forgeries, that manifest is refreshed intentionally; hashes prove byte identity, not whether the bytes describe a legal experiment. The candidate additionally re-derives cross-artifact constraints.

The primary endpoint is macro fault-detection recall. Success required recall ≥ 0.90 , no rejection of a clean trace, and an absolute recall gain ≥ 0.25 . The benchmark used local CPU only and accessed no scientific labels, protected




Figure 1. The evaluated workflow binds proposal, execution, independent evaluation, decision, and evidence release. The semantic grader checks both artifact integrity and cross-stage constraints; agent prose and Tea Time reflection are excluded from evidence.




Figure 2. Fault-detection recall over 48 preregistered corrupt traces (four seed-level points per method; 12 families per seed). Bars show aggregate exact counts, and no uncertainty bars are drawn because the finite catalog is reported without population inference. Both methods accepted all four clean controls.

Table 1. Frozen fault-injection result.

Method	Detected	Recall	Clean rejected
Manifest hashes	4/48	0.083	0/4
Semantic grader	48/48	1.000	0/4

partitions, acquisition labels, or remote resources.

4. Results

The semantic grader detected 48/48 corruptions (macro and micro recall 1.00) and accepted 4/4 clean controls. The manifest-only integrity ablation detected 4/48 (macro and micro recall 0.083), for a preregistered macro-recall gain of 0.917 (Table 1); macro and micro values coincide because every family has four cases. Every threshold was met. No clean rejection was observed in four controls, which is not a precise estimate of a population false-positive rate. Because the faults were white-box and repository-defined, the result demonstrates coverage of this contract, not robustness to arbitrary attacks.

Real-MACE negative case. We also reused a hash-verified, policy-attested exact-commit replay on one NVIDIA L4. It used `mace-torch` 0.3.16, a candidate-only checkpoint, a bounded approved label view, one optimizer step per iteration, and an independent aggregate evaluator. For validation force-component MAE in $\text{eV \AA}^{-1}$, iteration 1 improved by 0.053%, below the frozen 0.1% threshold, and was rejected. Iteration 2 improved by 0.133% and passed the metric threshold, but both iterations had `rerun_metric_delta=1.0` and failed the exact-equality reproducibility constraint, so both were rejected. Here 1.0 is a binary mismatch indicator, not the physical magnitude of numerical drift. Exact equality is especially conservative because numerical reproducibility can vary across software releases, platforms, and CPU/GPU execution even with identical seeds (PyTorch Contributors, 2026). The completed trace remained `infrastructure_only` and `claim-ineligible`.

This case demonstrates the policy consequence: a metric-only rule would have promoted iteration 2, whereas the frozen multi-constraint rule retained the negative result. It does not show that the candidate was scientifically worse, because the current binary rerun endpoint cannot distinguish harmless floating-point variation from decision-relevant instability.

5. Limitations and Conclusion

Our fault operators were written against the grader contract, so perfect recall may overestimate coverage of unanticipated corruption. Four clean controls do not tightly bound false-positive probability. The mutable manifest-only integrity ablation is deliberately narrow, and the experiment evaluates trace validity rather than scientific truth. The MACE replay contains two infrastructure iterations on one device and no frozen-test release, physical-property validation, independent provider replication, or approved checkpoint promotion. Repeated validation access can also create adaptive bias (Dwork et al., 2015); this campaign makes no claim about that risk.

Within these limits, the experiment supports one conclusion: for the twelve preregistered fault families, semantic re-derivation caught failures that producer-declared hashes usually could not, without rejecting the clean controls. In the accompanying MACE case, the same fail-closed posture preserved a nominal metric improvement as a rejection rather than converting it into an unsupported performance claim. Future work should replace the binary rerun flag with tolerance-aware, decision-calibrated deltas and evaluate externally designed faults.

References

Batatia, I., Kovacs, D. P., Simm, G., Ortner, C., and Csanyi, G. MACE: Higher order equivariant message passing neural networks for fast and accurate force fields. In *Advances in Neural Information Processing Systems 35*, pp. 11423–11436, 2022. doi: 10.52202/068431-0830. URL <https://doi.org/10.52202/068431-0830>.

Bouthillier, X., Laurent, C., and Vincent, P. Unreproducible research is reproducible. In *Proceedings of the 36th International Conference on Machine Learning*, volume 97 of *Proceedings of Machine Learning Research*, pp. 725–734. PMLR, 2019. URL <https://proceedings.mlr.press/v97/bouthillier19a.html>.

Dwork, C., Feldman, V., Hardt, M., Pitassi, T., Reingold, O., and Roth, A. The reusable holdout: Preserving validity in adaptive data analysis. *Science*, 349(6248):636–638, 2015. doi: 10.1126/science.aaa9375. URL <https://doi.org/10.1126/science.aaa9375>.

Hsueh, M.-C., Tsai, T. K., and Iyer, R. K. Fault injection techniques and tools. *Computer*, 30(4):75–82, 1997. doi: 10.1109/2.585157. URL <https://doi.org/10.1109/2.585157>.

Huang, K., Jin, Y., Li, R., Li, M. Y., Candes, E., and Leskovec, J. Automated hypothesis validation with agentic sequential falsifications. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pp. 25372–25437. PMLR, 2025. URL <https://proceedings.mlr.press/v267/huang25n.html>.

Huang, Q., Vora, J., Liang, P., and Leskovec, J. MLAGent-Bench: Evaluating language agents on machine learning experimentation. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pp. 20271–20309. PMLR, 2024. URL <https://proceedings.mlr.press/v235/huang24y.html>.

Maxson, T., Soyemi, A., Chen, B. W. J., and Szilvási, T. Enhancing the quality and reliability of machine

learning interatomic potentials through better reporting practices. *The Journal of Physical Chemistry C*, 128(16):6524–6537, 2024. doi: 10.1021/acs.jpcc.4c00028. URL <https://doi.org/10.1021/acs.jpcc.4c00028>.

PyTorch Contributors. Reproducibility. PyTorch 2.13 documentation, 2026. URL <https://docs.pytorch.org/docs/2.13/notes/randomness.html>. Accessed 2026-07-12.

Starace, G., Jaffe, O., Sherburn, D., Aung, J., Chan, J. S., Maksin, L., Dias, R., Mays, E., Kinsella, B., Thompson, W., Heidecke, J., Glaese, A., and Patwardhan, T. PaperBench: Evaluating AI’s ability to replicate AI research. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pp. 56843–56873. PMLR, 2025. URL <https://proceedings.mlr.press/v267/starace25a.html>.

Yamada, Y., Lange, R. T., Lu, C., Hu, S., Lu, C., Forster, J., Clune, J., and Ha, D. The AI scientist-v2: Workshop-level automated scientific discovery via agentic tree search, 2025. URL <https://arxiv.org/abs/2504.08066>.